

Spring Cloud Data Flow Samples

1.2.0.BUILD-SNAPSHOT



Table of Contents

1.	Overvi	ew	. 1
2.	Stream	ning	. 2
	2.1.	Cassandra Samples	. 2
		HTTP to Cassandra Demo	. 2
		Prerequisites	. 2
		Using the Local Server	2
		Using Cloud Foundry Server	. 4
		Summary	. 7
	2.2.	JDBC Samples	. 8
		HTTP to MySQL Demo	. 8
		Prerequisites	. 8
		Using Local Server	. 8
		Using Cloud Foundry Server	10
		Summary	13
	2.3.	GemFire Samples	13
		HTTP to Gemfire Demo	13
		Prerequisites	14
		Using the Local Server	14
		Using the Cloud Foundry Server	16
		Summary	19
		Gemfire to Log Demo	20
		Prerequisites	20
		Using the Local Server	21
		Using the Cloud Foundry Server	23
		Summary	26
		Gemfire CQ to Log Demo	26
		Prerequisites	26
		Using the Local Server	27
		Using the Cloud Foundry Server	29
		Summary	32
	2.4.	Custom Stream Application Samples	33
		Custom Spring Cloud Stream Processor	33
		Prerequisites	33
		Create the custom stream app	34
		Deploy the app to Spring Cloud Data Flow	35
		Summary	36
3.	Task /	Batch	37
	3.1.	Task Samples	37
		Batch Job on Cloud Foundry	37
		Prerequisites	37
		Summary	41
4.	Analyti	CS	42
	4.1.	Twitter Analytics	42
		Prerequisites	42
		Summary	
5.	Functio	ons	47
	5.1.	Functions in Spring Cloud Data Flow	47

Spring Cloud Data Flow Samples

Prerequisites	47
Summary	5 Ω

1. Overview

This guide contains samples and demonstrations of how to build data microservices applications with Spring Cloud Data Flow.

2. Streaming

2.1 Cassandra Samples

HTTP to Cassandra Demo

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from an *HTTP* endpoint and write the payload to a *Cassandra* database.

We will take you through the steps to configure and Spring Cloud Data Flow server in either a <u>local</u> or <u>Cloud Foundry</u> environment.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

Using the Local Server

Additional Prerequisites

A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-local/target

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- Running instance of Kafka
- · Running instance of Apache Cassandra
- A database utility tool such as <u>DBeaver</u> to connect to the Cassandra instance. You might have to provide host, port, username and password depending on the Cassandra configuration you are using.
- Create a keyspace and a book table in Cassandra using:

1. Register the out-of-the-box applications for the Kafka binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, repository, app import --uri bit.ly/Bacon-RELEASE-streamapplications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-kafka-10-maven
```

2. Create the stream

```
dataflow:>stream create cassandrastream --definition "http --server.port=8888 --
spring.cloud.stream.bindings.output.contentType='application/json' | cassandra --
ingestQuery='insert into book (id, isbn, title, author) values (uuid(), ?, ?, ?)' --
keyspace=clouddata" --deploy

Created and deployed new stream 'cassandrastream'
```



Note

If Cassandra isn't running on default port on localhost or if you need username and password to connect, use one of the following options to specify the necessary connection parameters: --username='<USERNAME>' --password='<PASSWORD>' --port=<PORT> --contact-points=<LIST-OF-HOSTS>

3. Verify the stream is successfully deployed

```
dataflow:>stream list
```

4. Notice that cassandrastream-http and cassandrastream-cassandra <u>Spring Cloud Stream</u> applications are running as Spring Boot applications within the server as a collocated process.

```
2015-12-15 15:52:31.576 INFO 18337 --- [nio-9393-exec-1]
o.s.c.d.a.s.l.OutOfProcessModuleDeployer : deploying module
org.springframework.cloud.stream.module:cassandra-sink:jar:exec:1.0.0.BUILD-SNAPSHOT instance 0
Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-data-
flow-284240942697761420/cassandrastream.cassandra
2015-12-15 15:52:31.583 INFO 18337 --- [nio-9393-exec-1]
o.s.c.d.a.s.l.OutOfProcessModuleDeployer : deploying module
org.springframework.cloud.stream.module:http-source:jar:exec:1.0.0.BUILD-SNAPSHOT instance 0
Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-data-
flow-284240942697761420/cassandrastream.http
```

5. Post sample data pointing to the http endpoint: localhost:8888 (8888 is the server.port we specified for the http source in this case)

```
dataflow:>http post --contentType 'application/json' --data '{"isbn": "1599869772", "title": "The
Art of War", "author": "Sun Tzu"}' --target http://localhost:8888
> POST (application/json;charset=UTF-8) http://localhost:8888 {"isbn": "1599869772", "title": "The
Art of War", "author": "Sun Tzu"}
> 202 ACCEPTED
```

6. Connect to the Cassandra instance and query the table clouddata.book to list the persisted records

```
select * from clouddata.book;
```

7. You're done!

Using Cloud Foundry Server

Additional Prerequisites

- · Cloud Foundry instance
- The Spring Cloud Data Flow Cloud Foundry Server

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-cloudfoundry/target



Note

Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

```
Welcome to the Spring Cloud Data Flow shell. For assistance hit TAB or type "help".
server-unknown:>
server-unknown:>dataflow config server http://dataflow-server.app.io
Successfully targeted http://dataflow-server.app.io
dataflow:>
```

- Running instance of cassandra in Cloud Foundry or from another Cloud provider
- A database utility tool such as <u>DBeaver</u> to connect to the Cassandra instance. You might have to provide host, port, username and password depending on the Cassandra configuration you are using.
- Create a book table in your Cassandra keyspace using:

```
CREATE TABLE book (

id uuid PRIMARY KEY,

isbn text,

author text,

title text
);
```

6. Register the out-of-the-box applications for the Rabbit binder



Note

address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

7. Create the stream

```
dataflow:>stream create cassandrastream --definition "http --
spring.cloud.stream.bindings.output.contentType='application/json' | cassandra --ingestQuery='insert
into book (id, isbn, title, author) values (uuid(), ?, ?, ?)' --username='<USERNAME>' --
password='<PASSWORD>' --port=<PORT> --contact-points=<HOST> --keyspace='<KEYSPACE>'" --deploy
Created and deployed new stream 'cassandrastream'
```

8. Verify the stream is successfully deployed

```
dataflow:>stream list
```

9. Notice that cassandrastream-http and cassandrastream-cassandra Spring Cloud Stream applications are running as *cloud-native* (microservice) applications in Cloud Foundry

10Lookup the url for cassandrastream-http application from the list above. Post sample data pointing to the http endpoint: <YOUR-cassandrastream-http-APP-URL>

```
http post --contentType 'application/json' --data '{"isbn": "1599869772", "title": "The Art of War", "author": "Sun Tzu"}' --target http://<YOUR-cassandrastream-http-APP-URL>
> POST (application/json;charset=UTF-8) http://cassandrastream-http.app.io {"isbn": "1599869772", "title": "The Art of War", "author": "Sun Tzu"}
> 202 ACCEPTED
```

11 Connect to the Cassandra instance and query the table book to list the data inserted

```
select * from book;
```

12Now, let's try to take advantage of Pivotal Cloud Foundry's platform capability. Let's scale the cassandrastream-http application from 1 to 3 instances

```
$ cf scale cassandrastream-http -i 3
Scaling app cassandrastream-http in org user-dataflow / space development as user...
OK
```

13. Verify App instances (3/3) running successfully

```
$ cf apps
Getting apps in org user-dataflow / space development as user...
OK
                     requested state instances memory disk urls
name
                                    1/1
cassandrastream-cassandra started
                                             1G 1G cassandrastream-
cassandra.app.io
                    started
                                 3/3 1G 1G cassandrastream-http.app.io
cassandrastream-http
                                    1/1
                                              1G
                                                     1G dataflow-server.app.io
dataflow-server
                     started
```

14.You're done!

Summary

In this sample, you have learned:

- How to use Spring Cloud Data Flow's Local and Cloud Foundry servers
- How to use Spring Cloud Data Flow's shell
- How to create streaming data pipeline to connect and write to Cassandra
- How to scale data microservice applications on Pivotal Cloud Foundry

2.2 JDBC Samples

HTTP to MySQL Demo

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from an http endpoint and write to MySQL database using jdbc sink.

We will take you through the steps to configure and Spring Cloud Data Flow server in either a <u>local</u> or <u>Cloud Foundry</u> environment.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

Using Local Server

Additional Prerequisites

A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-local/target</code>

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- Running instance of <u>Kafka</u>
- Running instance of MySQL
- A database utility tool such as <u>DBeaver</u> or <u>DbVisualizer</u>
- Create the test database with a names table (in MySQL) using:

```
CREATE DATABASE test;
USE test;
CREATE TABLE names
(
name varchar(255)
);
```

1. Register the out-of-the-box applications for the Kafka binder

corresponding to your installed app.



Note

These samples assume that the Data Flow Server can access a remote Maven repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow:>app import --uri bit.ly/Bacon-RELEASE-streamapplications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual

apps using dataflow:>app register... using the maven:// resource URI format

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-kafka-10-maven
```

2. Create the stream

```
dataflow:>stream create --name mysqlstream --definition "http --server.port=8787 | jdbc --
tableName=names --columns=name --spring.datasource.driver-class-name=org.mariadb.jdbc.Driver --
spring.datasource.url='jdbc:mysql://localhost:3306/test'" --deploy

Created and deployed new stream 'mysqlstream'
```



Note

If MySQL isn't running on default port on localhost or if you need username and password to connect, use one of the following options to specify the

```
necessary connection parameters: --spring.datasource.url='jdbc:mysql://
<HOST>:<PORT>/<NAME>' --spring.datasource.username=<USERNAME> --
spring.datasource.password=<PASSWORD>
```

3. Verify the stream is successfully deployed

```
dataflow:>stream list
```

4. Notice that mysqlstream-http and mysqlstream-jdbc Spring Cloud Stream applications are running as Spring Boot applications within the Local server as collocated processes.

```
2016-05-03 09:29:55.918 INFO 65162 --- [nio-9393-exec-3] o.s.c.d.spi.local.LocalAppDeployer
: deploying app mysqlstream.jdbc instance 0
Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-
dataflow-6850863945840320040/mysqlstream1-1462292995903/mysqlstream.jdbc
2016-05-03 09:29:55.939 INFO 65162 --- [nio-9393-exec-3] o.s.c.d.spi.local.LocalAppDeployer
: deploying app mysqlstream.http instance 0
Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-
dataflow-6850863945840320040/mysqlstream-1462292995934/mysqlstream.http
```

5. Post sample data pointing to the http endpoint: localhost:8787 [8787 is the server.port we specified for the http source in this case]

```
dataflow:>http post --contentType 'application/json' --target http://localhost:8787 --data "{\"name
\": \"Foo\"}"
> POST (application/json;charset=UTF-8) http://localhost:8787 {"name": "Spring Boot"}
> 202 ACCEPTED
```

+

1. Connect to the MySQL instance and guery the table test.names to list the new rows:

```
select * from test.names;
```

2. You're done!

Using Cloud Foundry Server

Additional Prerequisites

- Cloud Foundry instance
- The Spring Cloud Data Flow Cloud Foundry Server

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-cloudfoundry/target</code>



Note

Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
```

```
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

- Running instance of rabbit in Cloud Foundry
- Running instance of mysql in Cloud Foundry
- A database utility tool such as <u>DBeaver</u> or <u>DbVisualizer</u>
- Create the names table (in MySQL) using:

```
CREATE TABLE names
```

```
(
  name varchar(255)
);
```

6. Register the out-of-the-box applications for the Rabbit binder



Note

repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow: >app import --uri bit.ly/Bacon-RELEASE-stream-applicationsrabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

These samples assume that the Data Flow Server can access a remote Maven

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

7. Create the stream

```
dataflow:>stream create --name mysqlstream --definition "http | jdbc --tableName=names --
columns=name"
Created new stream 'mysqlstream'

dataflow:>stream deploy --name mysqlstream --properties
   "app.jdbc.spring.cloud.deployer.cloudfoundry.services=mysql"
Deployed stream 'mysqlstream'
```



Note

By supplying the app.jdbc.spring.cloud.deployer.cloudfoundry.services=mysql property, we are deploying the stream with jdbc-sink to automatically bind to mysql service and only this application in the stream gets the service binding. This also eliminates the requirement to supply datasource credentials in stream definition.

8. Verify the stream is successfully deployed

```
dataflow:>stream list
```

9. Notice that mysqlstream-http and mysqlstream-jdbc Spring Cloud Stream applications are running as *cloud-native* (microservice) applications in Cloud Foundry

```
$ cf apps
```

```
Getting apps in org user-dataflow / space development as user...

OK

name requested state instances memory disk urls

mysqlstream-http started 1/1 1G 1G mysqlstream-http.app.io

mysqlstream-jdbc started 1/1 1G 1G mysqlstream-jdbc.app.io

dataflow-server started 1/1 1G 1G dataflow-server.app.io
```

10Lookup the url for mysqlstream-http application from the list above. Post sample data pointing to the http endpoint: <YOUR-mysqlstream-http-APP-URL>

```
http post --contentType 'application/json' --target http://mysqlstream-http.app.io --data "{\"name\":
    \"Bar"}"
> POST (application/json; charset=UTF-8) http://mysqlstream-http.app.io {"name": "Bar"}
> 202 ACCEPTED
```

11 Connect to the MySQL instance and query the table names to list the new rows:

```
select * from names;
```

12Now, let's take advantage of Pivotal Cloud Foundry's platform capability. Let's scale the mysqlstream-http application from 1 to 3 instances

```
$ cf scale mysqlstream-http -i 3
Scaling app mysqlstream-http in org user-dataflow / space development as user...
OK
```

13. Verify App instances (3/3) running successfully

14.You're done!

Summary

In this sample, you have learned:

- How to use Spring Cloud Data Flow's Local and Cloud Foundry servers
- How to use Spring Cloud Data Flow's shell
- How to create streaming data pipeline to connect and write to MySQL
- How to scale data microservice applications on Pivotal Cloud Foundry

2.3 GemFire Samples

HTTP to Gemfire Demo

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from an http endpoint and write to Gemfire using the gemfire sink.

We will take you through the steps to configure and run Spring Cloud Data Flow server in either a <u>local</u> or <u>Cloud Foundry</u> environment.



Note

For legacy reasons the <code>gemfire</code> Spring Cloud Stream Apps are named after <code>PivotalGemFire</code>. The code base for the commercial product has since been open sourced as <code>ApacheGeode</code>. These samples should work with compatible versions of Pivotal GemFire or Apache Geode. Herein we will refer to the installed IMDG simply as <code>Geode</code>.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for <u>download</u> or you can <u>build</u> it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI <u>localhost:9393/dashboard</u>, (or wherever it the server is hosted) to perform the same operations.

 A Geode installation with a locator and cache server running Unresolved directive in streaming/ gemfire/gemfire-http/overview.adoc - include::geode-setup.adoc[]

Using the Local Server

Additional Prerequisites

A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-local/target</code>

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

A running instance of <u>Rabbit MQ</u>

1. Use gfsh to start a locator and server

```
gfsh>start locator --name=locator1
gfsh>start server --name=server1
```

2. Create a region called Stocks

```
gfsh>create region --name Stocks --type=REPLICATE
```

Use the Shell to create the sample stream

1. Register the out-of-the-box applications for the Rabbit binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow: >app import --uri bit.ly/Bacon-Release-stream-applications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, <a href="source-http=maven://org.springframework.cloud.stream.app:http-source-rabbit:1.2.0.Release registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The

address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

2. Create the stream

This example creates an http endpoint to which we will post stock prices as a JSON document containing symbol and price fields. The property --json=true to enable Geode's JSON support and configures the sink to convert JSON String payloads to PdxInstance, the recommended way to store JSON documents in Geode. The keyExpression property is a SpEL expression used to extract the symbol value the PdxInstance to use as an entry key.



Note

PDX serialization is very efficient and supports OQL queries without requiring a custom domain class. Use of custom domain types requires these classes to be in the class path of both the stream apps and the cache server. For this reason, the use of custom payload types is generally discouraged.

```
dataflow:>stream create --name stocks --definition "http --port=9090 | gemfire --json=true -- regionName=Stocks --keyExpression=payload.getField('symbol')" --deploy
Created and deployed new stream 'stocks'
```



Note

If the Geode locator isn't running on default port on localhost, add the options -- connect-type=locator --host-addresses=<host>:<port>. If there are multiple locators, you can provide a comma separated list of locator addresses. This is not necessary for the sample but is typical for production environments to enable fail-over.

3. Verify the stream is successfully deployed

```
dataflow:>stream list
```

4. Post sample data pointing to the http endpoint: localhost:9090 (9090 is the port we specified for the http source)

```
dataflow:>http post --target http://localhost:9090 --contentType application/json --data
  '{"symbol":"VMW","price":117.06}'
> POST (application/json) http://localhost:9090 {"symbol":"VMW","price":117.06}
> 202 ACCEPTED
```

5. Using gfsh, connect to the locator if not already connected, and verify the cache entry was created.

```
gfsh>get --key='VMW' --region=/Stocks
Result : true
Key Class : java.lang.String
Key : VMW
Value Class : org.apache.geode.pdx.internal.PdxInstanceImpl

symbol | price
----- | -----
VMW | 117.06
```

6. You're done!

Using the Cloud Foundry Server

Additional Prerequisites

- · A Cloud Foundry instance
- · Cloud Data Flow Cloud Foundry Server

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-cloudfoundry/target</code>



Note

Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

Successfully targeted http://dataflow-server.app.io

dataflow:>

- Running instance of a rabbit service in Cloud Foundry
- Running instance of the <u>Pivotal Cloud Cache for PCF</u> (PCC) service cloudcache in Cloud Foundry.
- 6. Register the out-of-the-box applications for the Rabbit binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow: >app import --uri bit.ly/Bacon-Release-stream-applications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, <a href="source-http=maven://org.springframework.cloud.stream.app:http-source-rabbit:1.2.0.Release registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The

address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

7. Get the PCC connection information

```
$ cf service-key cloudcache my-service-key
Getting key my-service-key for service instance cloudcache as <user>...
 "locators": [
 "10.0.16.9[55221]",
 "10.0.16.11[55221]",
  "10.0.16.10[55221]"
 "urls": {
 "gfsh": "http://...",
  "pulse": "http://.../pulse"
 },
 "users": [
 {
   "password": <password>,
  "username": "cluster_operator"
   "password": <password>,
   "username": "developer"
]
```

8. Using gfsh, connect to the PCC instance as cluster_operator using the service key values and create the Stocks region.

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>create region --name Stocks --type=REPLICATE
```

9. Create the stream, connecting to the PCC instance as developer

This example creates an http endpoint to which we will post stock prices as a JSON document containing symbol and price fields. The property --json=true to enable Geode's JSON support and configures the sink to convert JSON String payloads to PdxInstance, the recommended way to store JSON documents in Geode. The keyExpression property is a SpEL expression used to extract the symbol value the PdxInstance to use as an entry key.



Note

PDX serialization is very efficient and supports OQL queries without requiring a custom domain class. Use of custom domain types requires these classes to be in the class path of both the stream apps and the cache server. For this reason, the use of custom payload types is generally discouraged.

```
dataflow:>stream create --name stocks --definition "http --security.basic.enabled=false | gemfire --username=developer --password=<developer-password> --connect-type=locator --host-addresses=10.0.16.9:55221 --regionName=Stocks --keyExpression=payload.getField('symbol')" --deploy
```

10.Verify the stream is successfully deployed

```
dataflow:>stream list
```

11Post sample data pointing to the http endpoint

Get the url of the http source using cf apps

```
dataflow:>http post --target http://<http source url> --contentType application/json --data
'{"symbol":"VMW","price":117.06}'
> POST (application/json) http://... {"symbol":"VMW","price":117.06}
> 202 ACCEPTED
```

12Using gfsh, connect to the PCC instance as cluster operator using the service key values.

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>get --key='VMW' --region=/Stocks
Result : true
Key Class : java.lang.String
Key : VMW
Value Class : org.apache.geode.pdx.internal.PdxInstanceImpl

symbol | price
----- | ------
VMW | 117.06
```

13.You're done!

Summary

In this sample, you have learned:

• How to use Spring Cloud Data Flow's Local and Cloud Foundry servers

- How to use Spring Cloud Data Flow's shell
- How to create streaming data pipeline to connect and write to gemfire

Gemfire to Log Demo

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from a <code>gemfire</code> endpoint and write to a log using the <code>log</code> sink. The <code>gemfire</code> source creates a <u>CacheListener</u> to monitor events for a region and publish a message whenever an entry is changed.

We will take you through the steps to configure and run Spring Cloud Data Flow server in either a <u>local</u> or Cloud Foundry environment.



Note

For legacy reasons the <code>gemfire</code> Spring Cloud Stream Apps are named after <code>Pivotal GemFire</code>. The code base for the commercial product has since been open sourced as <code>Apache Geode</code>. These samples should work with compatible versions of Pivotal GemFire or Apache Geode. Herein we will refer to the installed IMDG simply as <code>Geode</code>.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

 A Geode installation with a locator and cache server running Unresolved directive in streaming/ gemfire/gemfire-log/overview.adoc - include::geode-setup.adoc[]

Using the Local Server

Additional Prerequisites

A Running Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-local/target

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

A running instance of Rabbit MQ

1. Use gfsh to start a locator and server

```
gfsh>start locator --name=locator1
gfsh>start server --name=server1
```

2. Create a region called Test

```
gfsh>create region --name Test --type=REPLICATE
```

Use the Shell to create the sample stream.

1. Register the out-of-the-box applications for the Rabbit binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, <a href="repositorgootnote-server-se

rabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://sgroupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven://resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

2. Create the stream

This example creates an gemfire source to which will publish events on a region

```
dataflow:>stream create --name events --definition " gemfire --regionName=Test | log" --deploy Created and deployed new stream 'events'
```



Note

If the Geode locator isn't running on default port on localhost, add the options --connect-type=locator --host-addresses=<host>:<port>. If there are multiple locators, you can provide a comma separated list of locator addresses. This is not necessary for the sample but is typical for production environments to enable fail-over.

3. Verify the stream is successfully deployed

```
dataflow:>stream list
```

4. Monitor stdout for the log sink. When you deploy the stream, you will see log messages in the Data Flow server console like this

```
2017-10-28 17:28:23.275 INFO 15603 --- [nio-9393-exec-2] o.s.c.d.spi.local.LocalAppDeployer :
Deploying app with deploymentId events.log instance 0.
Logs will be in /var/folders/hd/5yqz2v2d3sxd3n879f4sg4gr0000gn/T/spring-cloud-
dataflow-4093992067314402881/events-1509226103269/events.log
2017-10-28 17:28:23.277 INFO 15603 --- [nio-9393-exec-2] o.s.c.d.s.c.StreamDeploymentController
: Downloading resource URI [maven://org.springframework.cloud.stream.app:gemfire-source-
rabbit:1.2.0.RELEASE]
2017-10-28 17:28:23.311 INFO 15603 --- [nio-9393-exec-2] o.s.c.d.s.c.StreamDeploymentController :
Deploying application named [gemfire] as part of stream named [events] with resource URI [maven://
org.springframework.cloud.stream.app:gemfire-source-rabbit:1.2.0.RELEASE]
2017-10-28 17:28:23.318 INFO 15603 --- [nio-9393-exec-2] o.s.c.d.spi.local.LocalAppDeployer :
Deploying app with deploymentId events.gemfire instance 0.
Logs will be in /var/folders/hd/5yqz2v2d3sxd3n879f4sg4gr0000gn/T/spring-cloud-
dataflow-4093992067314402881/events-1509226103311/events.gemfire
```

Copy the location of the log sink logs. This is a directory that ends in events.log. The log files will be in stdout_0.log under this directory. You can monitor the output of the log sink using tail, or something similar:

```
\label{thm:condition} $$ tail -f /var/folders/hd/5yqz2v2d3sxd3n879f4sg4gr0000gn/T/spring-cloud-dataflow-4093992067314402881/events-1509226103269/events.log/stdout_0.log
```

5. Using gfsh, create and update some cache entries

```
gfsh>put --region /Test --key 1 --value "value 1"
gfsh>put --region /Test --key 2 --value "value 2"
gfsh>put --region /Test --key 3 --value "value 3"
gfsh>put --region /Test --key 1 --value "new value 1"
```

6. Observe the log output You should see messages like:

```
2017-10-28 17:28:52.893 INFO 18986 --- [emfire.events-1] log sink : value 1"
2017-10-28 17:28:52.893 INFO 18986 --- [emfire.events-1] log sink : value 2"
2017-10-28 17:28:52.893 INFO 18986 --- [emfire.events-1] log sink : value 3"
```

```
2017-10-28 17:28:52.893 INFO 18986 --- [emfire.events-1] log sink : new value 1"
```

By default, the message payload contains the updated value. Depending on your application, you may need additional information. The data comes from EntryEvent. You can access any fields using the source's cache-event-expression property. This takes a SpEL expression bound to the EntryEvent. Try something like --cache-event-expression '{key:'+key+',new_value:'+newValue+'}' (HINT: You will need to destroy the stream and recreate it to add this property, an exercise left to the reader). Now you should see log messages like:

```
2017-10-28 17:28:52.893 INFO 18986 --- [emfire.events-1] log-sink : {key:1,new_value:value 1} 2017-10-28 17:41:24.466 INFO 18986 --- [emfire.events-1] log-sink : {key:2,new_value:value 2}
```

7. You're done!

Using the Cloud Foundry Server

Additional Prerequisites

- · A Cloud Foundry instance
- The Spring Cloud Data Flow Cloud Foundry Server

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-cloudfoundry/target



Note

Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

- Running instance of a rabbit service in Cloud Foundry
- Running instance of the <u>Pivotal Cloud Cache for PCF</u> (PCC) service cloudcache in Cloud Foundry.
- 6. Register the out-of-the-box applications for the Rabbit binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, <a href="repository.com/re

address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using

whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven://resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

7. Get the PCC connection information

```
$ cf service-key cloudcache my-service-key
Getting key my-service-key for service instance cloudcache as <user>...
 "locators": [
  "10.0.16.9[55221]".
 "10.0.16.11[55221]",
  "10.0.16.10[55221]"
 "urls": {
 "gfsh": "http://...",
  "pulse": "http://.../pulse"
 "users": [
   "password": <password>,
  "username": "cluster_operator"
   "password": <password>,
  "username": "developer"
 ]
}
```

8. Using gfsh, connect to the PCC instance as cluster_operator using the service key values and create the Test region.

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>create region --name Test --type=REPLICATE
```

9. Create the stream, connecting to the PCC instance as developer. This example creates an gemfire source to which will publish events on a region

```
dataflow stream create --name events --definition " gemfire --username=developer -- password=<developer-password> --connect-type=locator --host-addresses=10.0.16.9:55221 -- regionName=Test | log" --deploy
```

10. Verify the stream is successfully deployed

```
dataflow:>stream list
```

11 Monitor stdout for the log sink

```
cf logs <log-sink-app-name>
```

12.Using gfsh, create and update some cache entries

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>put --region /Test --key 1 --value "value 1"
gfsh>put --region /Test --key 2 --value "value 2"
gfsh>put --region /Test --key 3 --value "value 3"
gfsh>put --region /Test --key 1 --value "new value 1"
```

13.Observe the log output

You should see messages like:

By default, the message payload contains the updated value. Depending on your application, you may need additional information. The data comes from EntryEvent. You can access any fields using the source's cache-event-expression property. This takes a SpEL expression bound to the EntryEvent. Try something like (HINT: You will need to destroy the stream and recreate it to add this property, an exercise left to the reader). Now you should see log messages like:

14.You're done!

Summary

In this sample, you have learned:

- How to use Spring Cloud Data Flow's Local and Cloud Foundry servers
- How to use Spring Cloud Data Flow's shell
- · How to create streaming data pipeline to connect and publish events from gemfire

Gemfire CQ to Log Demo

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from a <code>gemfire-cq</code> (Continuous Query) endpoint and write to a log using the <code>log</code> sink. The <code>gemfire-cq</code> source creates a Continuous Query to monitor events for a region that match the query's result set and publish a message whenever such an event is emitted. In this example, we simulate monitoring orders to trigger a process whenever the quantity ordered is above a defined limit.

We will take you through the steps to configure and run Spring Cloud Data Flow server in either a <u>local</u> or Cloud Foundry environment.



Note

For legacy reasons the <code>gemfire</code> Spring Cloud Stream Apps are named after <code>Pivotal</code> <code>GemFire</code>. The code base for the commercial product has since been open sourced as <code>Apache</code> <code>Geode</code>. These samples should work with compatible versions of Pivotal GemFire or Apache Geode. Herein we will refer to the installed IMDG simply as <code>Geode</code>.

Prerequisites

A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI <u>localhost:9393/dashboard</u>, (or wherever it the server is hosted) to perform the same operations.

 A Geode installation with a locator and cache server running Unresolved directive in streaming/ gemfire/gemfire-cq-log/overview.adoc - include::geode-setup.adoc[]

Using the Local Server

Additional Prerequisites

A Running Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-local/target

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- A running instance of <u>Rabbit MQ</u>
- 1. Use gfsh to start a locator and server

```
gfsh>start locator --name=locator1
gfsh>start server --name=server1
```

2. Create a region called Orders

```
gfsh>create region --name Orders --type=REPLICATE
```

==== Use the Shell to create the sample stream.

3. Register the out-of-the-box applications for the Rabbit binder



Note

repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow: > app import --uri bit.ly/Bacon-RELEASE-stream-applicationsrabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven://resource URI format corresponding to

These samples assume that the Data Flow Server can access a remote Maven

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

4. Create the stream

your installed app.

This example creates an gemfire-cq source to which will publish events matching a query criteria on a region. In this case we will monitor the Orders region. For simplicity, we will avoid creating a data structure for the order. Each cache entry contains an integer value representing the quantity of the ordered item. This stream will fire a message whenever the value>999. By default, the source emits only the value. Here we will override that using the cq-event-expression property. This accepts a SpEL expression bound to a CQEvent. To reference the entire CQEvent instace, we use #this. In order to display the contents in the log, we will invoke toString() on the instance.

```
dataflow:>stream create --name orders --definition " gemfire-cq --query='SELECT * from /Orders o where o > 999' --cq-event-expression=#this.toString() | log" --deploy Created and deployed new stream 'events'
```



Note

If the Geode locator isn't running on default port on localhost, add the options -- connect-type=locator --host-addresses=<host>:<port>. If there are multiple

locators, you can provide a comma separated list of locator addresses. This is not necessary for the sample but is typical for production environments to enable fail-over.

5. Verify the stream is successfully deployed

```
dataflow:>stream list
```

6. Monitor stdout for the log sink. When you deploy the stream, you will see log messages in the Data Flow server console like this

```
2017-10-30 09:39:36.283 INFO 8167 --- [nio-9393-exec-5] o.s.c.d.spi.local.LocalAppDeployer :
Deploying app with deploymentId orders.log instance 0.
Logs will be in /var/folders/hd/5yqz2v2d3sxd3n879f4sg4gr0000gn/T/spring-cloud-dataflow-5375107584795488581/orders-1509370775940/orders.log
```

Copy the location of the log sink logs. This is a directory that ends in orders.log. The log files will be in stdout_0.log under this directory. You can monitor the output of the log sink using tail, or something similar:

```
\label{thm:condition} $$ tail -f /var/folders/hd/5yqz2v2d3sxd3n879f4sg4gr0000gn/T/spring-cloud-dataflow-5375107584795488581/orders-1509370775940/orders.log/stdout_0.log
```

7. Using gfsh, create and update some cache entries

```
gfsh>put --region Orders --value-class java.lang.Integer --key 01234 --value 1000
gfsh>put --region Orders --value-class java.lang.Integer --key 11234 --value 1005
gfsh>put --region Orders --value-class java.lang.Integer --key 21234 --value 100
gfsh>put --region Orders --value-class java.lang.Integer --key 31234 --value 999
gfsh>put --region Orders --value-class java.lang.Integer --key 21234 --value 1000
```

8. Observe the log output You should see messages like:

```
2017-10-30 09:53:02.231 INFO 8563 --- [ire-cq.orders-1] log-sink :

CqEvent [CqName=GfCq1; base operation=CREATE; cq operation=CREATE; key=01234; value=1000]

2017-10-30 09:53:19.732 INFO 8563 --- [ire-cq.orders-1] log-sink :

CqEvent [CqName=GfCq1; base operation=CREATE; cq operation=CREATE; key=11234; value=1005]

2017-10-30 09:53:53.242 INFO 8563 --- [ire-cq.orders-1] log-sink :

CqEvent [CqName=GfCq1; base operation=UPDATE; cq operation=CREATE; key=21234; value=1000]
```

9. Another interesting demonstration combines gemfire-cg with the http-qemfire example.

```
dataflow:> stream create --name stocks --definition "http --port=9090 | gemfire-json-server -- regionName=Stocks --keyExpression=payload.getField('symbol')" --deploy dataflow:> stream create --name stock_watch --definition "gemfire-cq --query='Select * from /Stocks where symbol=''VMW''' | log" --deploy
```

1. You're done!

Using the Cloud Foundry Server

Additional Prerequisites

- · A Cloud Foundry instance
- The Spring Cloud Data Flow Cloud Foundry Server

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-cloudfoundry/target</code>



Note

Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

Successfully targeted http://dataflow-server.app.io

dataflow:>

- Running instance of a rabbit service in Cloud Foundry
- Running instance of the <u>Pivotal Cloud Cache for PCF</u> (PCC) service cloudcache in Cloud Foundry.
- 6. Register the out-of-the-box applications for the Rabbit binder



Note

These samples assume that the Data Flow Server can access a remote Maven repository, <a href="repositorgo:reposito

address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

7. Get the PCC connection information

```
$ cf service-key cloudcache my-service-key
Getting key my-service-key for service instance cloudcache as <user>...
 "locators": [
 "10.0.16.9[55221]",
 "10.0.16.11[55221]",
  "10.0.16.10[55221]"
 "urls": {
 "gfsh": "http://...",
  "pulse": "http://.../pulse"
 },
 "users": [
 {
   "password": <password>,
  "username": "cluster_operator"
   "password": <password>,
   "username": "developer"
]
```

8. Using gfsh, connect to the PCC instance as cluster_operator using the service key values and create the Test region.

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>create region --name Orders --type=REPLICATE
```

9. Create the stream using the Data Flow Shell

This example creates an gemfire-cq source to which will publish events matching a query criteria on a region. In this case we will monitor the Orders region. For simplicity, we will avoid creating a data structure for the order. Each cache entry contains an integer value representing the quantity of the ordered item. This stream will fire a message whenever the value>999. By default, the source emits only the value. Here we will override that using the cq-event-expression property. This accepts a SpEL expression bound to a CQEvent. To reference the entire CQEvent instance, we use #this. In order to display the contents in the log, we will invoke toString() on the instance.

```
dataflow:>stream create --name orders --definition "gemfire-cq --username=developer --
password=<developer-password> --connect-type=locator --host-addresses=10.0.16.9:55221 --query='SELECT
* from /Orders o where o > 999' --cq-event-expression=#this.toString() | log" --deploy
Created and deployed new stream 'events'
```

10.Verify the stream is successfully deployed

```
dataflow:>stream list
```

11 Monitor stdout for the log sink

```
cf logs <log-sink-app-name>
```

12Using gfsh, create and update some cache entries

```
gfsh>connect --use-http --url=<gfsh-url> --user=cluster_operator --
password=<cluster_operator_password>
gfsh>put --region Orders --value-class java.lang.Integer --key 01234 --value 1000
gfsh>put --region Orders --value-class java.lang.Integer --key 11234 --value 1005
gfsh>put --region Orders --value-class java.lang.Integer --key 21234 --value 100
gfsh>put --region Orders --value-class java.lang.Integer --key 31234 --value 999
gfsh>put --region Orders --value-class java.lang.Integer --key 21234 --value 1000
```

13. Observe the log output You should see messages like:

```
2017-10-30 09:53:02.231 INFO 8563 --- [ire-cq.orders-1] log-sink :
CqEvent [CqName=GfCq1; base operation=CREATE; cq operation=CREATE; key=01234; value=1000]
2017-10-30 09:53:19.732 INFO 8563 --- [ire-cq.orders-1] log-sink :
CqEvent [CqName=GfCq1; base operation=CREATE; cq operation=CREATE; key=11234; value=1005]
2017-10-30 09:53:53.242 INFO 8563 --- [ire-cq.orders-1] log-sink :
CqEvent [CqName=GfCq1; base operation=UPDATE; cq operation=CREATE; key=21234; value=1000]
```

14Another interesting demonstration combines gemfire-cq with the http-gemfire example.

```
dataflow:> stream create --name stocks --definition "http --port=9090 | gemfire-json-server -- regionName=Stocks --keyExpression=payload.getField('symbol')" --deploy dataflow:> stream create --name stock_watch --definition "gemfire-cq --query='Select * from /Stocks where symbol=''VMW''' | log" --deploy
```

1. You're done!

Summary

In this sample, you have learned:

- How to use Spring Cloud Data Flow's Local and Cloud Foundry servers
- How to use Spring Cloud Data Flow's shell
- How to create streaming data pipeline to connect and publish CQ events from gemfire

2.4 Custom Stream Application Samples

Custom Spring Cloud Stream Processor

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-local/target</code>

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- A Java IDE
- Maven Installed
- · A running instance of Rabbit MQ

Create the custom stream app

We will create a custom <u>Spring Cloud Stream</u> application and run it on Spring Cloud Data Flow. We'll go through the steps to make a simple processor that converts temperature from Fahrenheit to Celsius. We will be running the demo locally, but all the steps will work in a Cloud Foundry environment as well.

- 1. Create a new spring cloud stream project
 - Create a <u>Spring initializer</u> project
 - Set the group to demo.celsius.converter and the artifact name as celsius-converterprocessor
 - Choose a message transport binding as a dependency for the custom app There are options for choosing Rabbit MQ or Kafka as the message transport. For this demo, we will use rabbit. Type rabbit in the search bar under Search for dependencies and select Stream Rabbit.
 - Hit the generate project button and open the new project in an IDE of your choice

2. Develop the app

We can now create our custom app. Our Spring Cloud Stream application is a Spring Boot application that runs as an executable jar. The application will include two Java classes:

- CelsiusConverterProcessorAplication.java the main Spring Boot application class, generated by Spring initializr
- CelsiusConverterProcessorConfiguration.java the Spring Cloud Stream code that we will write

We are creating a transformer that takes a Fahrenheit input and converts it to Celsius. Following the same naming convention as the application file, create a new Java class in the same package called CelsiusConverterProcessorConfiguration.java.

CelsiusConverterProcessorConfiguration.java.

```
@EnableBinding(Processor.class)
public class CelsiusConverterProcessorConfiguration {

    @Transformer(inputChannel = Processor.INPUT, outputChannel = Processor.OUTPUT)
    public int convertToCelsius(String payload) {
        int fahrenheitTemperature = Integer.parseInt(payload);
        return (farenheitTemperature-32)*5/9;
    }
}
```

Here we introduced two important Spring annotations. First we annotated the class with <code>@EnableBinding(Processor.class)</code>. Second we created a method and annotated it with <code>@Transformer(inputChannel = Processor.INPUT)</code>, outputChannel = <code>Processor.OUTPUT)</code>. By adding these two annotations we have configured this stream app as a <code>Processor</code> (as opposed to a <code>Source</code> or a <code>Sink</code>). This means that the application receives

input from an upstream application via the Processor.INPUT channel and sends its output to a downstream application via the Processor.OUTPUT channel.

The convertToCelsius method takes a String as input for Fahrenheit and then returns the converted Celsius as an integer. This method is very simple, but that is also the beauty of this programming style. We can add as much logic as we want to this method to enrich this processor. As long as we annotate it properly and return valid output, it works as a Spring Cloud Stream Processor. Also note that it is straightforward to unit test this code.

3. Build the Spring Boot application with Maven

```
$cd <PROJECT_DIR>
$./mvnw clean package
```

4. Run the Application standalone

```
java -jar target/celsius-converter-processor-0.0.1-SNAPSHOT.jar
```

If all goes well, we should have a running standalone Spring Boot Application. Once we verify that the app is started and running without any errors, we can stop it.

Deploy the app to Spring Cloud Data Flow

1. Register the out-of-the-box applications for the Rabbit binder



Note

rabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://sgroupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow:>app register... using the maven:// resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

2. Register the custom processor

```
app register --type processor --name convertToCelsius --uri <File URL of the jar file on the local filesystem where you built the project above> --force
```

3. Create the stream

We will create a stream that uses the out of the box http source and log sink and our custom transformer.

```
dataflow:>stream create --name convertToCelsiusStream --definition "http --port=9090 | convertToCelsius | log" --deploy

Created and deployed new stream 'convertToCelsiusStream'
```

4. Verify the stream is successfully deployed

```
dataflow:>stream list
```

5. Verify that the apps have successfully deployed

```
dataflow:>runtime apps

2016-09-27 10:03:11.988 INFO 95234 --- [nio-9393-exec-9] o.s.c.d.spi.local.LocalAppDeployer : deploying app convertToCelsiusStream.log instance 0
Logs will be in /var/folders/2q/krqwcbhj2d58csmthyq_n1nw0000gp/T/spring-cloud-dataflow-323689888473815319/convertToCelsiusStream-1474984991968/convertToCelsiusStream.log 2016-09-27 10:03:12.397 INFO 95234 --- [nio-9393-exec-9] o.s.c.d.spi.local.LocalAppDeployer : deploying app convertToCelsiusStream.convertToCelsius instance 0
Logs will be in /var/folders/2q/krqwcbhj2d58csmthyq_n1nw0000gp/T/spring-cloud-dataflow-3236898888473815319/convertToCelsiusStream-1474984992392/convertToCelsiusStream.convertToCelsius
2016-09-27 10:03:14.445 INFO 95234 --- [nio-9393-exec-9] o.s.c.d.spi.local.LocalAppDeployer : deploying app convertToCelsiusStream.http instance 0
Logs will be in /var/folders/2q/krqwcbhj2d58csmthyq_n1nw0000gp/T/spring-cloud-dataflow-3236898888473815319/convertToCelsiusStream-1474984994440/convertToCelsiusStream.http
```

6. Post sample data to the http endpoint: localhost:9090 (9090 is the port we specified for the http source in this case)

```
dataflow:>http post --target http://localhost:9090 --data 76
> POST (text/plain;Charset=UTF-8) http://localhost:9090 76
> 202 ACCEPTED
```

7. Open the log file for the convertToCelsiusStream.log app to see the output of our stream

```
tail -f /var/folders/2q/krqwcbhj2d58csmthyq_n1nw0000gp/T/spring-cloud-dataflow-7563139704229890655/convertToCelsiusStream-1474990317406/convertToCelsiusStream.log/stdout_0.log
```

You should see the temperature you posted converted to Celsius!

```
2016-09-27 10:05:34.933 INFO 95616 --- [CelsiusStream-1] log.sink : 24
```

Summary

- How to write a custom Processor stream application
- How to use Spring Cloud Data Flow's Local server
- How to use Spring Cloud Data Flow's shell application

3. Task / Batch

3.1 Task Samples

Batch Job on Cloud Foundry

In this demonstration, you will learn how to orchestrate short-lived data processing application (*eg:* Spring Batch Jobs) using Spring Cloud Task and Spring Cloud Data Flow on Cloud Foundry.

Prerequisites

- Local PCFDev instance
- Local install of cf CLI command line tool
- · Running instance of mysgl in PCFDev
- A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

The Spring Cloud Data Flow Cloud Foundry Server running in PCFDev

The Cloud Foundry Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-cloudfoundry/target</code>



Although you can run the Data Flow Cloud Foundry Server locally and configure it to deploy to any Cloud Foundry instance, we will deploy the server to Cloud Foundry as recommended.

1. Verify that CF instance is reachable (Your endpoint urls will be different from what is shown here).

```
$ cf api
API endpoint: https://api.system.io (API version: ...)
$ cf apps
Getting apps in org [your-org] / space [your-space] as user...
OK
No apps found
```

2. Follow the instructions to deploy the <u>Spring Cloud Data Flow Cloud Foundry server</u>. Don't worry about creating a Redis service. We won't need it. If you are familiar with Cloud Foundry application manifests, we recommend creating a manifest for the the Data Flow server as shown <u>here</u>.



Warning

As of this writing, there is a typo on the SPRING_APPLICATION_JSON entry in the sample manifest. SPRING_APPLICATION_JSON must be followed by : and The JSON string must be wrapped in single quotes. Alternatively, you can replace that line with MAVEN_REMOTE_REPOSITORIES_REPO1_URL: repo.spring.io/libs-snapshot. If your Cloud Foundry installation is behind a firewall, you may need to install the stream apps used in this sample in your internal Maven repository and configure the server to access that repository.

3. Once you have successfully executed cf push, verify the dataflow server is running

- 4. Notice that the dataflow-server application is started and ready for interaction via the url endpoint
- 5. Connect the shell with server running on Cloud Foundry, e.g., dataflow-server.app.io

Successfully targeted http://dataflow-server.app.io



Note

PCF 1.7.12 or greater is required to run Tasks on Spring Cloud Data Flow. As of this writing, PCFDev and PWS supports builds upon this version.

6. Task support needs to be enabled on pcf-dev. Being logged as admin, issue the following command:

```
cf enable-feature-flag task_creation
Setting status of task_creation as admin...

OK

Feature task_creation Enabled.
```



Note

For this sample, all you need is the <code>mysql</code> service and in PCFDev, the <code>mysql</code> service comes with a different plan. From CF CLI, create the service by: <code>cf</code> <code>create-service</code> <code>p-mysql</code> 512mb <code>mysql</code> and bind this service to <code>dataflow-server</code> by: <code>cf</code> <code>bind-service</code> <code>dataflow-server</code> <code>mysql</code>.



Note

All the apps deployed to PCFDev start with low memory by default. It is recommended to change it to at least 768MB for dataflow-server. Ditto for every app spawned by Spring Cloud Data Flow. Change the memory by: cf setenv dataflow-server SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_MEMORY 512. Likewise, we would have to skip SSL validation by: cf set-env dataflow-server SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_SKIP_SSL_VALIDATION true.

7. Tasks in Spring Cloud Data Flow require an RDBMS to host "task repository" (see here for more details), so let's instruct the Spring Cloud Data Flow server to bind the mysql service to each deployed task:

```
$ cf set-env dataflow-server SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_TASK_SERVICES mysql
$ cf restage dataflow-server
```



Note

We only need mysql service for this sample.

8. As a recap, here is what you should see as configuration for the Spring Cloud Data Flow server:

```
Cf env dataflow-server

....

User-Provided:

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_DOMAIN: local.pcfdev.io

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_MEMORY: 512

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_ORG: pcfdev-org

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_PASSWORD: pass

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_SKIP_SSL_VALIDATION: false

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_SPACE: pcfdev-space

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_TASK_SERVICES: mysql

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_URL: https://api.local.pcfdev.io

SPRING_CLOUD_DEPLOYER_CLOUDFOUNDRY_USERNAME: user
```

```
No running env variables have been set

No staging env variables have been set
```

- 9. Notice that dataflow-server application is started and ready for interaction via <u>dataflow-server.local.pcfdev.io</u> endpoint
- 10Build and register the batch-job <u>example</u> from Spring Cloud Task samples. For convenience, the final <u>uber-jar artifact</u> is provided with this sample.

```
dataflow:>app register --type task --name simple_batch_job --uri https://github.com/spring-cloud/spring-cloud-dataflow-samples/raw/master/tasks/simple-batch-job/batch-job-1.3.0.BUILD-SNAPSHOT.jar
```

11 Create the task with simple-batch-job application

```
dataflow:>task create foo --definition "simple_batch_job"
```



Note

Unlike Streams, the Task definitions don't require explicit deployment. They can be launched on-demand, scheduled, or triggered by streams.

12. Verify there's **still** no Task applications running on PCFDev - they are listed only after the initial launch/staging attempt on PCF

13Let's launch foo

```
dataflow:>task launch foo
```

14. Verify the execution of foo by tailing the logs

```
2016-08-14T18:49:07.53-0700 [APP/TASK/foo/0]OUT 2016-08-15 01:49:07.536 INFO 14 --- [
main] o.s.b.c.l.support.SimpleJobLauncher : Job: [SimpleJob: [name=job2]] completed with the following parameters: [{}] and the following status: [COMPLETED]

...

2016-08-14T18:49:07.71-0700 [APP/TASK/foo/0]OUT Exit status 0
2016-08-14T18:49:07.78-0700 [APP/TASK/foo/0]OUT Destroying container
2016-08-14T18:49:08.47-0700 [APP/TASK/foo/0]OUT Successfully destroyed container
```



Verify job1 and job2 operations embedded in simple-batch-job application are launched independently and they returned with the status COMPLETED.



Note

Unlike LRPs in Cloud Foundry, tasks are short-lived, so the logs aren't always available. They are generated only when the Task application runs; at the end of Task operation, the container that ran the Task application is destroyed to free-up resources.

15List Tasks in Cloud Foundry

```
$ cf apps
Getting apps in org pcfdev-org / space pcfdev-space as user...
OK

name requested state instances memory disk urls
dataflow-server started 1/1 768M 512M dataflow-server.local.pcfdev.io
foo stopped 0/1 1G 1G
```

16.Verify Task execution details

17. Verify Job execution details

Summary

- · How to register and orchestrate Spring Batch jobs in Spring Cloud Data Flow
- How to use the cf CLI in the context of Task applications orchestrated by Spring Cloud Data Flow
- How to verify task executions and task repository

4. Analytics

4.1 Twitter Analytics

In this demonstration, you will learn how to build a data pipeline using <u>Spring Cloud Data Flow</u> to consume data from *TwitterStream* and compute simple analytics over data-in-transit using *Field-Value-Counter*.

We will take you through the steps to configure Spring Cloud Data Flow's Local server.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

· A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in spring-cloud-dataflow-server-local/target

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- · Running instance of Redis
- Running instance of <u>Kafka</u>
- Twitter credentials from Twitter Developers site
 - 1. Register the out-of-the-box applications for the Kafka binder



These samples assume that the Data Flow Server can access a remote Maven repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow:>app import --uri bit.ly/Bacon-RELEASE-streamapplications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow: >app register... using the maven: // resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-kafka-10-maven
```

2. Create and deploy the following streams

```
(1) dataflow:>stream create tweets --definition "twitterstream --consumerKey=<CONSUMER_KEY>
--consumerSecret=<CONSUMER_SECRET> --accessToken=<ACCESS_TOKEN> --
accessTokenSecret=<ACCESS_TOKEN_SECRET> | log"
Created new stream 'tweets'

(2) dataflow:>stream create tweetlang --definition ":tweets.twitterstream > field-value-counter
--fieldName=lang --name=language" --deploy
Created and deployed new stream 'tweetlang'

(3) dataflow:>stream create tagcount --definition ":tweets.twitterstream > field-value-counter --
fieldName=entities.hashtags.text --name=hashtags" --deploy
Created and deployed new stream 'tagcount'

(4) dataflow:>stream deploy tweets
Deployed stream 'tweets'
```



Note

To get a consumerKey and consumerSecret you need to register a twitter application. If you don't already have one set up, you can create an app at the Twitter Developers site to get these credentials. The tokens Consumer_Secret, Consumer_Key, Consumer_Key), Consumer_Key, Consumer_Key), Consumer_Key), Consumer_Key), Consumer_Key), <a href="Consumer_Key), Consumer_Key), <a href="Consumer_Key), <a href="Consumer_Key</

3. Verify the streams are successfully deployed. Where: (1) is the primary pipeline; (2) and (3) are tapping the primary pipeline with the DSL syntax <stream-name>.<label/app name> [e.x.:tweets.twitterstream]; and (4) is the final deployment of primary pipeline

```
dataflow:>stream list
```

4. Notice that tweetlang.field-value-counter, tagcount.field-value-counter, tweets.log and tweets.twitterstream Spring Cloud Stream applications are running as Spring Boot applications within the local-server.

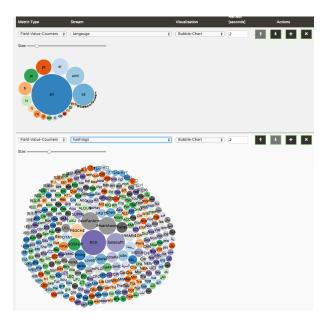
```
2016-02-16 11:43:26.174 INFO 10189 --- [nio-9393-exec-2] o.s.c.d.d.l.OutOfProcessModuleDeployer
       : deploying module org.springframework.cloud.stream.module:field-value-counter-
sink:jar:exec:1.0.0.BUILD-SNAPSHOT instance 0
      Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-data-
\verb|flow-6990537012958280418| tweetlang-1455651806160| tweetlang.field-value-counter| the counter of the counte
2016-02-16 11:43:26.206 INFO 10189 --- [nio-9393-exec-3] o.s.c.d.d.l.OutOfProcessModuleDeployer
       : deploying module org.springframework.cloud.stream.module:field-value-counter-
sink:jar:exec:1.0.0.BUILD-SNAPSHOT instance 0
    Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-data-
flow-6990537012958280418/tagcount-1455651806202/tagcount.field-value-counter
2016-02-16 11:43:26.806 INFO 10189 --- [nio-9393-exec-4] o.s.c.d.d.l.OutOfProcessModuleDeployer
      : deploying module org.springframework.cloud.stream.module:log-sink:jar:exec:1.0.0.BUILD-
SNAPSHOT instance 0
     Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gp/T/spring-cloud-data-
flow-6990537012958280418/tweets-1455651806800/tweets.log
2016-02-16 11:43:26.813 INFO 10189 --- [nio-9393-exec-4] o.s.c.d.d.l.OutOfProcessModuleDeployer
     : deploying module org.springframework.cloud.stream.module:twitterstream-
source:jar:exec:1.0.0.BUILD-SNAPSHOT instance 0
      Logs will be in /var/folders/c3/ctx7 rns6x30tg7rb76wzgwr0000gp/T/spring-cloud-data-
flow-6990537012958280418/tweets-1455651806800/tweets.twitterstream
```

5. Verify that two field-value-counter with the names hashtags and language is listing successfully

6. Verify you can query individual field-value-counter results successfully

```
dataflow:>field-value-counter display hashtags
Displaying values for field value counter 'hashtags'
Value
                              #Count#
#KCA
                             # 40#
#PENNYSTOCKS
                                17#
#TEAMBILLIONAIRE
                                17#
#UCL
                                 11#
# . . .
                                 ..#
                                ..#
#...
                              #
# . . .
                                 ..#
dataflow:>field-value-counter display language
Displaying values for field value counter 'language'
#############
#Value#Count#
#############
#en #1,171#
    # 337#
#ar # 296#
#und # 251#
#pt # 175#
#ja # 137#
#.. #
      ...#
#.. # ...#
    # ...#
# . .
#############
```

- 7. Go to Dashboard accessible at localhost:9393/dashboard and launch the Analytics tab. From the default Dashboard menu, select the following combinations to visualize real-time updates on field-value-counter.
 - For real-time updates on language tags, select:
 - a. Metric Type as Field-Value-Counters
 - b. Stream as language
 - c. Visualization as Bubble-Chart or Pie-Chart
 - For real-time updates on hashtags tags, select:
 - a. Metric Type as Field-Value-Counters
 - b. Stream as hashtags
 - c. Visualization as Bubble-Chart or Pie-Chart



Summary

- How to use Spring Cloud Data Flow's Local server
- How to use Spring Cloud Data Flow's shell application
- How to create streaming data pipeline to compute simple analytics using Twitter Stream and Field Value Counter applications

5. Functions

5.1 Functions in Spring Cloud Data Flow

In this sample, you will learn how to use <u>Spring Cloud Function</u> based streaming applications in Spring Cloud Data Flow. To learn more about Spring Cloud Function, check out the <u>project page</u>.

Prerequisites

· A Running Data Flow Shell

The Spring Cloud Data Flow Shell is available for download or you can build it yourself.



Note

the Spring Cloud Data Flow Shell and Local server implementation are in the same repository and are both built by running ./mvnw install from the project root directory. If you have already run the build, use the jar in spring-cloud-dataflow-shell/target

To run the Shell open a new terminal session:



Note

The Spring Cloud Data Flow Shell is a Spring Boot application that connects to the Data Flow Server's REST API and supports a DSL that simplifies the process of defining a stream or task and managing its lifecycle. Most of these samples use the shell. If you prefer, you can use the Data Flow UI localhost:9393/dashboard, (or wherever it the server is hosted) to perform the same operations.

· A running local Data Flow Server

The Local Data Flow Server is Spring Boot application available for <u>download</u> or you can <u>build</u> it yourself. If you build it yourself, the executable jar will be in <code>spring-cloud-dataflow-server-local/target</code>

To run the Local Data Flow server Open a new terminal session:

```
$cd <PATH/TO/SPRING-CLOUD-DATAFLOW-LOCAL-JAR>
$java -jar spring-cloud-dataflow-server-local-<VERSION>.jar
```

- A local build of Spring Cloud Function
- A running instance of Rabbit MQ
- General understanding of the out-of-the-box function-runner application
 - 1. Register the out-of-the-box applications for the Rabbit binder



These samples assume that the Data Flow Server can access a remote Maven repository, repo.spring.io/libs-release by default. If your Data Flow server is running behind a firewall, or you are using a maven proxy preventing access to public repositories, you will need to install the sample apps in your internal Maven repository and configure the server accordingly. The sample applications are typically registered using Data Flow's bulk import facility. For example, the Shell command dataflow:>app import --uri bit.ly/Bacon-RELEASE-streamapplications-rabbit-maven (The actual URI is release and binder specific so refer to the sample instructions for the actual URL). The bulk import URI references a plain text file containing entries for all of the publicly available Spring Cloud Stream and Task applications published to repo.spring.io. For example, source.http=maven:// org.springframework.cloud.stream.app:http-sourcerabbit:1.2.0.RELEASE registers the http source app at the corresponding Maven address, relative to the remote repository(ies) configured for the Data Flow server. The format is maven://<groupId>:<artifactId>:<version> You will need to download the required apps or build them and then install them in your Maven repository, using whatever group, artifact, and version you choose. If you do this, register individual apps using dataflow: >app register... using the maven: // resource URI format corresponding to your installed app.

```
dataflow:>app import --uri http://bit.ly/Bacon-RELEASE-stream-applications-rabbit-maven
```

2. Register the out-of-the-box <u>function-runner</u> application (we will use the 1.0.0.BUILD-SNAPSHOT built by the Spring CI system)

dataflow:>app register --name function-runner --type processor --uri http://repo.spring.io/libs-snapshot/org/springframework/cloud/stream/app/function-app-rabbit/1.0.0.BUILD-SNAPSHOT/function-app-rabbit-1.0.0.BUILD-SNAPSHOT.jar --metadata-uri http://repo.spring.io/libs-snapshot/org/springframework/cloud/stream/app/function-app-rabbit/1.0.0.BUILD-SNAPSHOT/function-app-rabbit-1.0.0.BUILD-SNAPSHOT-metadata.jar

3. Create and deploy the following stream

dataflow:>stream create foo --definition "http --server.port=9001 | function-runner -function.className=com.example.functions.CharCounter --function.location=file:///<PATH/TO/SPRINGCLOUD-FUNCTION>/spring-cloud-function-samples/function-sample/target/spring-cloud-functionsample-1.0.0.BUILD-SNAPSHOT.jar | log" --deploy



Note

Replace the <PATH/TO/SPRING-CLOUD-FUNCTION> with the correct path.



Note

The source core of CharCounter function is in Spring cloud Function's samples repo.

4. Verify the stream is successfully deployed.

```
dataflow:>stream list
#Stream#
Stream Definition
    # Status
# Name #
#foo #http --server.port=9001 | function-runner -
{\tt function.className=com.example.functions.CharCounter}
                            #All apps
   #--function.location=file:///<PATH/TO/SPRING-CLOUD-FUNCTION>/spring-cloud-function-samples/
function-sample/target/spring-cloud-function-sample-1.0.0.BUILD-<SNAPSHOT class="jar"></SNAPSHOT>
  #have been #
   #| log
   #successfully#
#
   #
   #deployed
```

Notice that foo-http, foo-function-runner, and foo-log <u>Spring Cloud Stream</u> applications
are running as Spring Boot applications and the log locations will be printed in the Local-server
console.

```
2017-10-17 11:43:03.714 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.s.s.AppDeployerStreamDeployer
  : Deploying application named [log] as part of stream named [foo] with resource URI [maven://
org.springframework.cloud.stream.app:log-sink-rabbit:jar:1.2.0.RELEASE]
2017-10-17 11:43:04.379 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.spi.local.LocalAppDeployer
  : Deploying app with deploymentId foo.log instance \ensuremath{\text{0}}.
  Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gs/T/spring-cloud-
dataflow-6549025456609489200/foo-1508265783715/foo.log
2017-10-17 11:43:04.380 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.s.s.AppDeployerStreamDeployer
    : Deploying application named [function-runner] as part of stream named [foo] with
resource URI [file:/var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gs/T/deployer-resource-
cache8941581850579153886/http-c73a62adae0abd7ec0dee91d891575709f02f8c9]
2017-10-17 11:43:04.384 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.spi.local.LocalAppDeployer
  : Deploying app with deploymentId foo.function-runner instance \ensuremath{\text{0}}.
  Logs will be in /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gs/T/spring-cloud-
dataflow-6549025456609489200/foo-1508265784380/foo.function-runner
2017-10-17 11:43:04.385 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.s.s.AppDeployerStreamDeployer
  : Deploying application named [http] as part of stream named [foo] with resource URI [maven://
org.springframework.cloud.stream.app:http-source-rabbit:jar:1.2.0.RELEASE]
2017-10-17 11:43:04.391 INFO 18409 --- [nio-9393-exec-7] o.s.c.d.spi.local.LocalAppDeployer
  : Deploying app with deploymentId foo.http instance 0.
  Logs will be in /var/folders/c3/ctx7 rns6x30tg7rb76wzgwr0000gs/T/spring-cloud-
dataflow-6549025456609489200/foo-1508265784385/foo.http
. . . .
```

6. Post sample data pointing to the http endpoint: localhost:9001 (9001 is the port we specified for the http source in this case)

```
dataflow:>http post --target http://localhost:9001 --data "hello world"
> POST (text/plain) http://localhost:9001 hello world
> 202 ACCEPTED

dataflow:>http post --target http://localhost:9001 --data "hmm, yeah, it works now!"
> POST (text/plain) http://localhost:9001 hmm, yeah, it works now!
> 202 ACCEPTED
```

7. Tail the log-sink's standard-out logs to see the character counts

```
$ tail -f /var/folders/c3/ctx7_rns6x30tq7rb76wzqwr0000gs/T/spring-cloud-
dataflow-6549025456609489200/foo-1508265783715/foo.log/stdout_0.log

....
....
2017-10-17 11:45:39.363 INFO 19193 --- [on-runner.foo-1] log-sink
: 11
2017-10-17 11:46:40.997 INFO 19193 --- [on-runner.foo-1] log-sink
: 24
....
....
```

Summary

- How to use Spring Cloud Data Flow's Local server
- How to use Spring Cloud Data Flow's shell application
- How to use the out-of-the-box function-runner application in Spring Cloud Data Flow